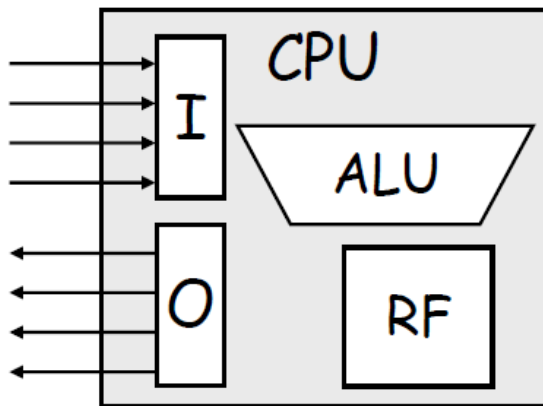


# اصول سیستم‌های کامپیوتری

---

## فصل نهم

### ورودی/خروجی - وقفه



# Copyright Notice

---

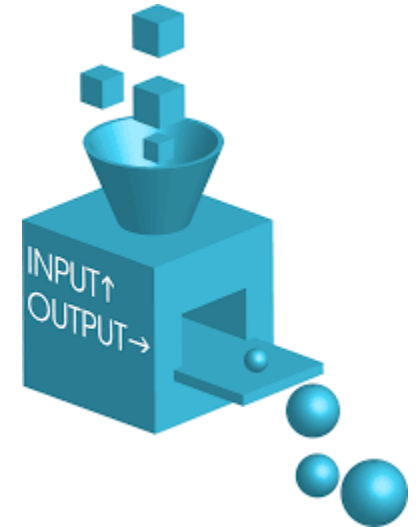
*Parts (text & figures) of this lecture are adopted from :*

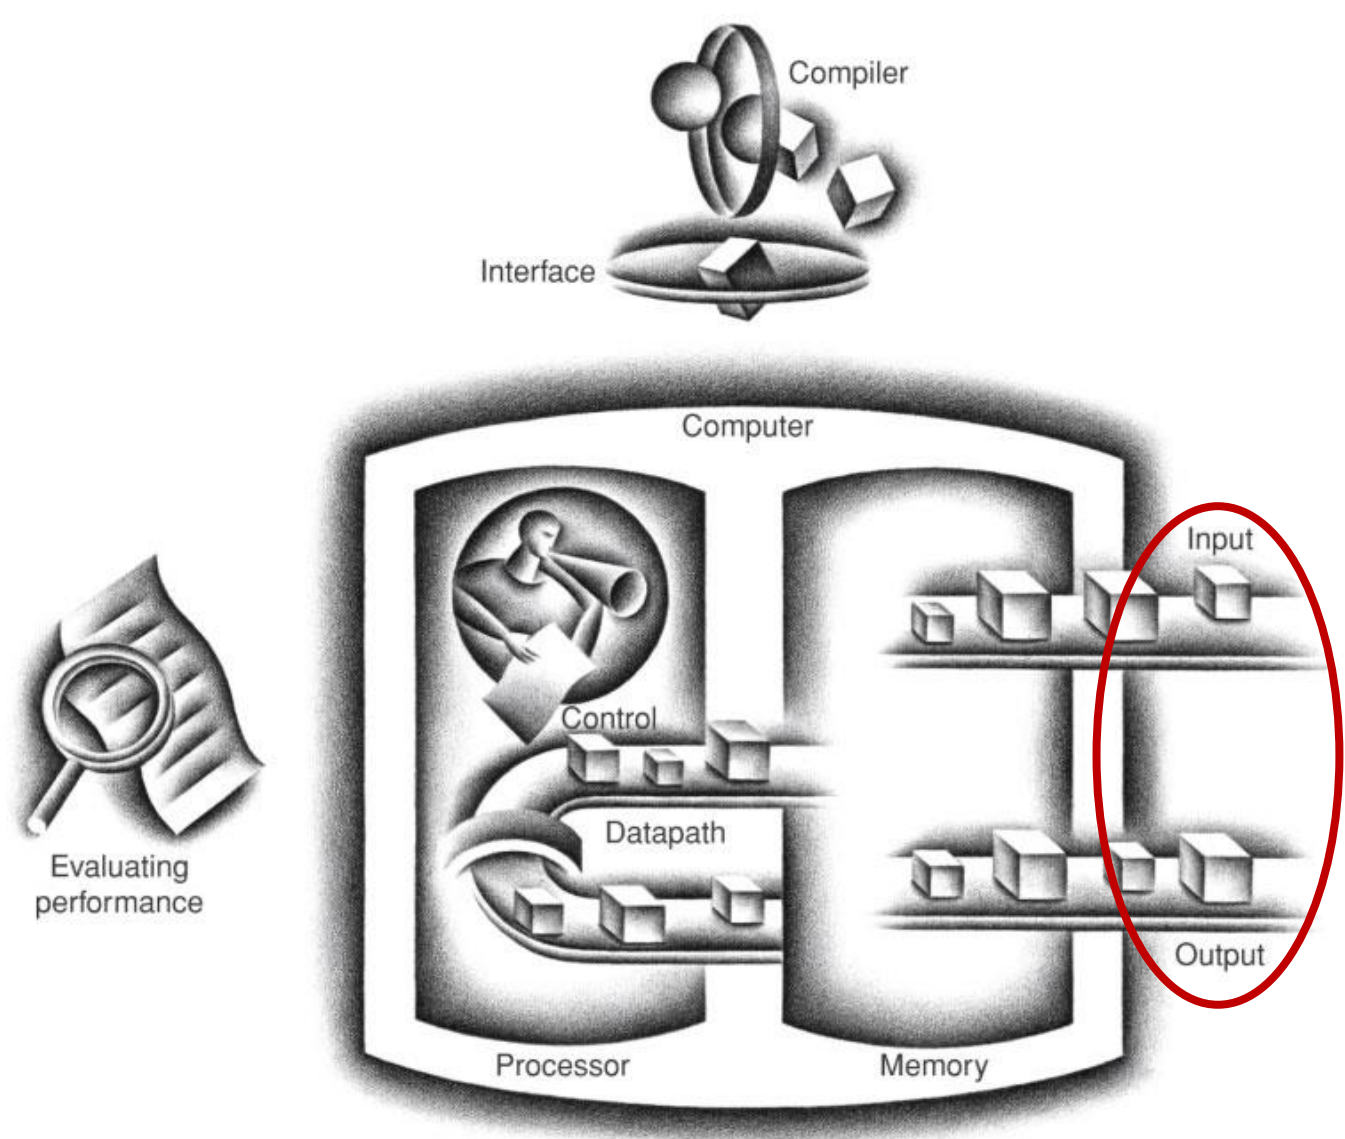
- 🔍 *D. Patterson & J. Hennessey, “Computer Organization & Design, The Hardware/Software Interface”, 5<sup>th</sup> Ed., MK publishing, 2014*
- 🔍 *M. Mano, C. Kime, T. Martin, “Logic and Computer Design Fundamentals”, 5<sup>th</sup> Ed., Pearson, 2015*
- 🔍 *W. Stallings, “Computer Organization and Architecture, Designing for Performance”, 10<sup>th</sup> Ed., Pearson, 2016*

# Outlines

---

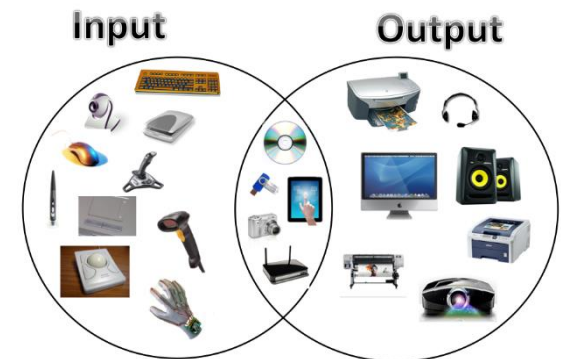
- *I/O devices (Instances & Characteristics)*
- *I/O transactions*
  - *Sync & Async data transfer*
  - *Handshaking*
- *ISA perspective*
  - *Memory-mapped vs. isolated I/O*
  - *Interrupt vs. Programmed I/O*
  - *Direct Memory Access (DMA)*





# I/O Devices

- Same components for all kinds of computer
  - Desktop, server, embedded
- Input/output includes
  - User-interface devices
    - Display, keyboard, mouse
  - Storage devices
    - Hard disk, CD/DVD, flash
  - Network adapters
    - For communicating with other computers



# *I/O Devices Characteristics - 1*

---

## ○ *Behavior:*

- *Input (read only)*
- *output (write only, cannot be read)*
- *storage (can be reread and usually rewritten)*

# I/O Devices Characteristics - 2

---

- Partner:

- A *human* or a *machine* is at the other end of the I/O device
  - feeding data on input or
  - reading data on output

# *I/O Devices Characteristics - 3*

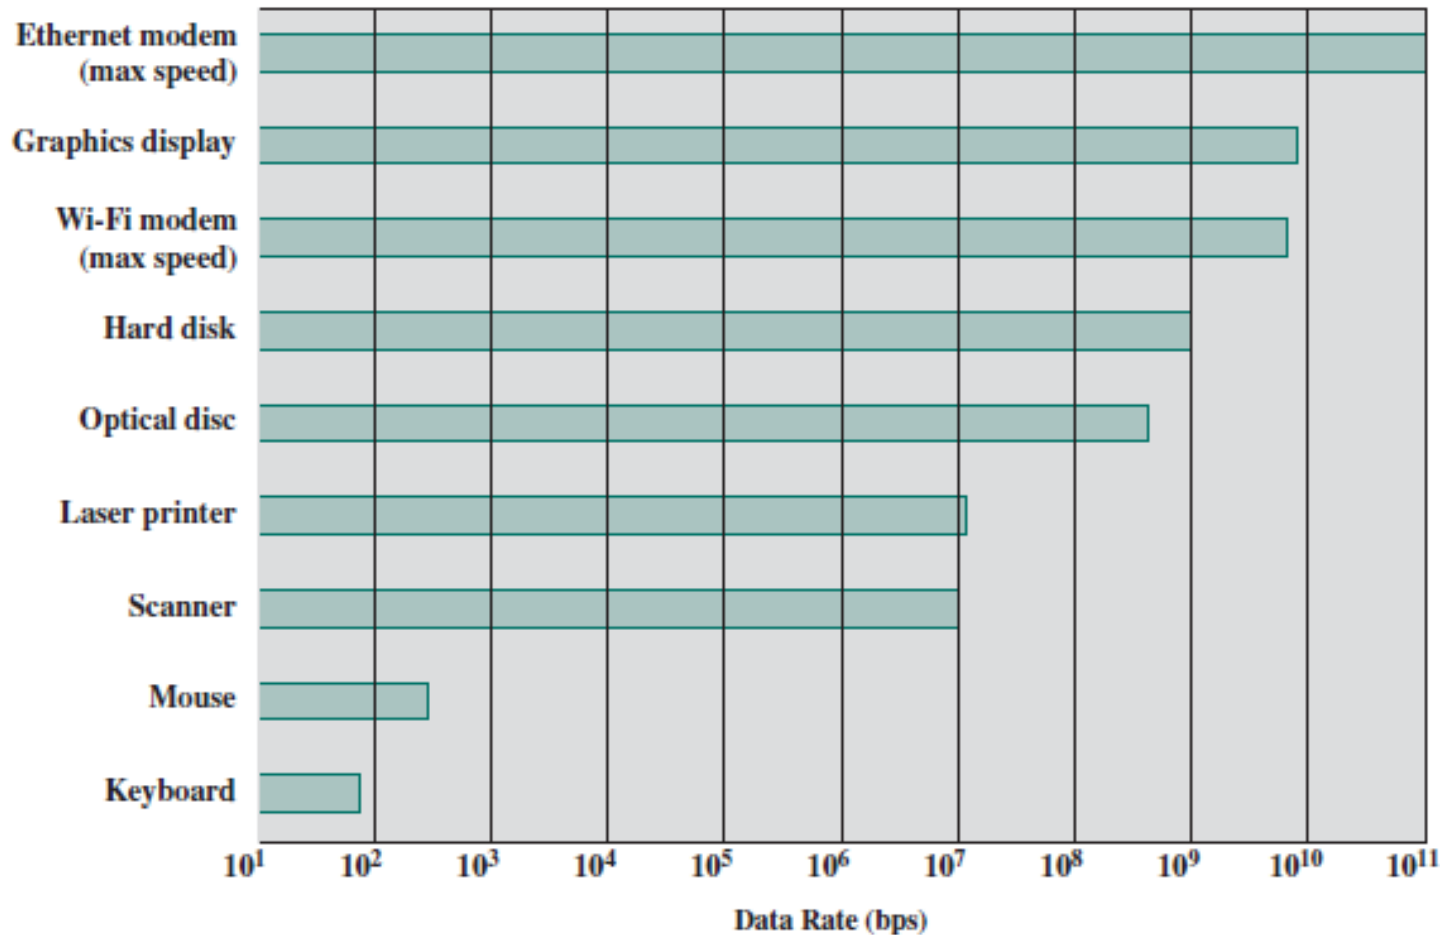
---

## ○ *Data rate:*

- *The peak rate at which data can be transferred between the I/O device & memory or processor*



# Typical I/O Device Data Rates



# *I/O Devices Characteristics*

---

- *Behavior (input, output, storage)*
- *Partner (a human or a machine)*
- *Data rate*
- *e.g. keyboard:*
  - *an input device*
  - *used by a human*
  - *with a peak data rate of about 10 bytes per second*



# Diversity of I/O Devices

Device	Behavior	Partner	Data rate (Mbit/sec)
Keyboard	Input	Human	0.0001
Mouse	Input	Human	0.0038
Voice input	Input	Human	0.2640
Sound input	Input	Machine	3.0000
Scanner	Input	Human	3.2000
Voice output	Output	Human	0.2640
Sound output	Output	Human	8.0000
Laser printer	Output	Human	3.2000
Graphics display	Output	Human	800.0000–8000.0000
Cable modem	Input or output	Machine	0.1280–6.0000
Network/LAN	Input or output	Machine	100.0000–10000.0000
Network/wireless LAN	Input or output	Machine	11.0000–54.0000
Optical disk	Storage	Machine	80.0000–220.0000
Magnetic tape	Storage	Machine	5.0000–120.0000
Flash memory	Storage	Machine	32.0000–200.0000
Magnetic disk	Storage	Machine	800.0000–3000.0000

# I/O Controller (I/O Bridge)

---

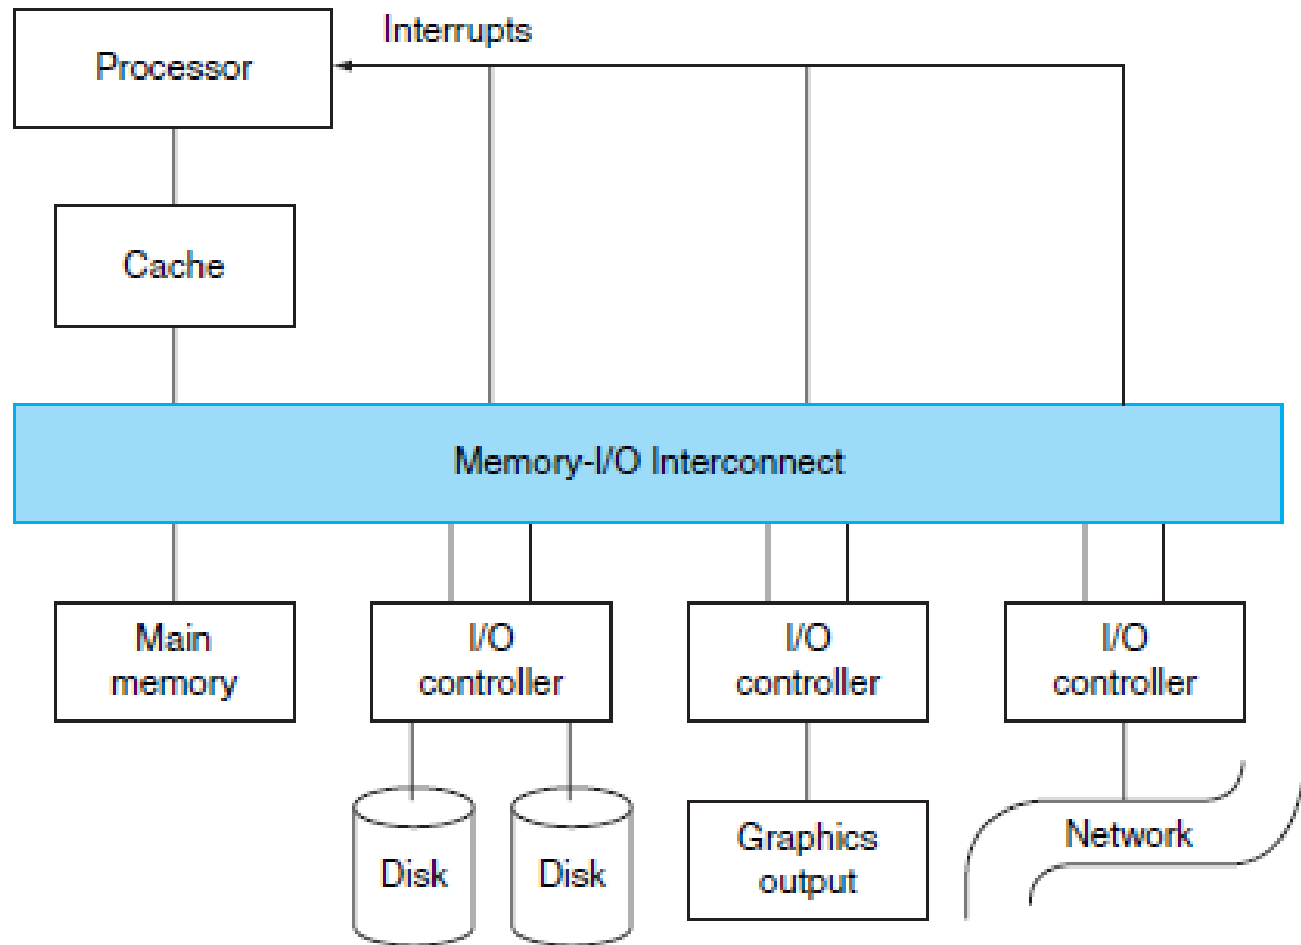
## ○ Fact

- *Input/outputs very slow devices*
  - *Average response time: few milliseconds*
- *CPUs very fast devices*
  - *Average running time: nanoseconds*

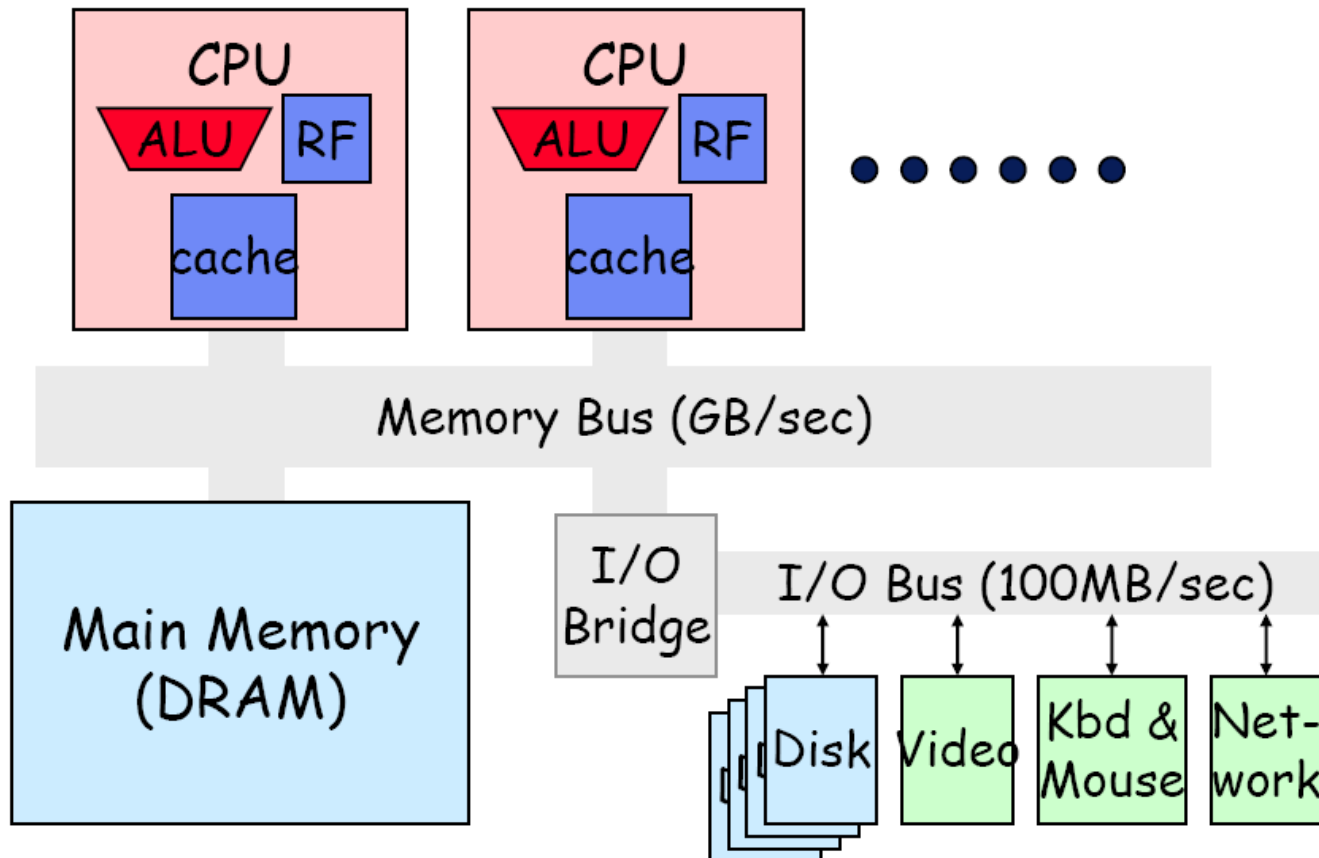
## ○ Question:

- *How are I/O devices connected to CPU?*
  - *I/O controller (also called I/O bridge)*

# Typical Collection of I/O Devices

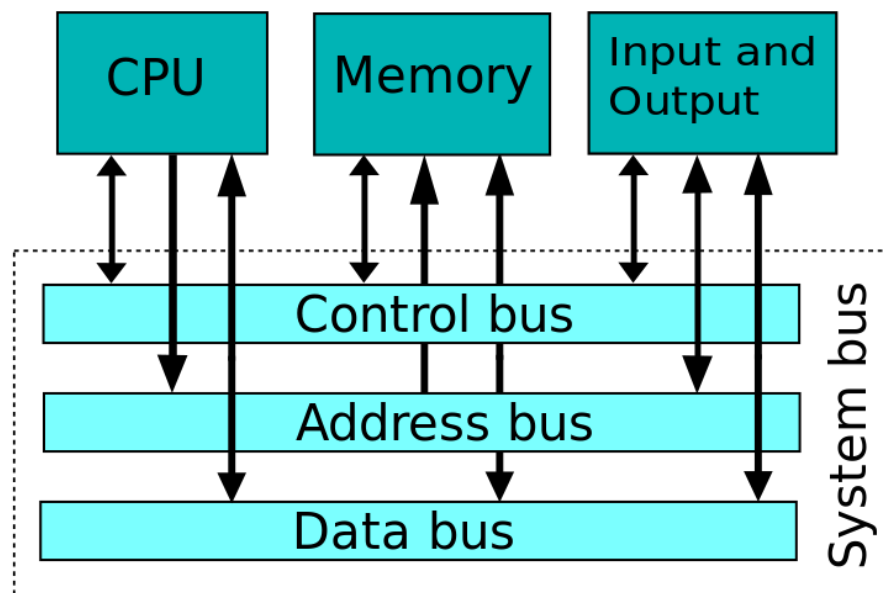


# Memory-I/O Separate Bus



# BUS

*Memory, CPU & I/O devices communicate via **BUS***



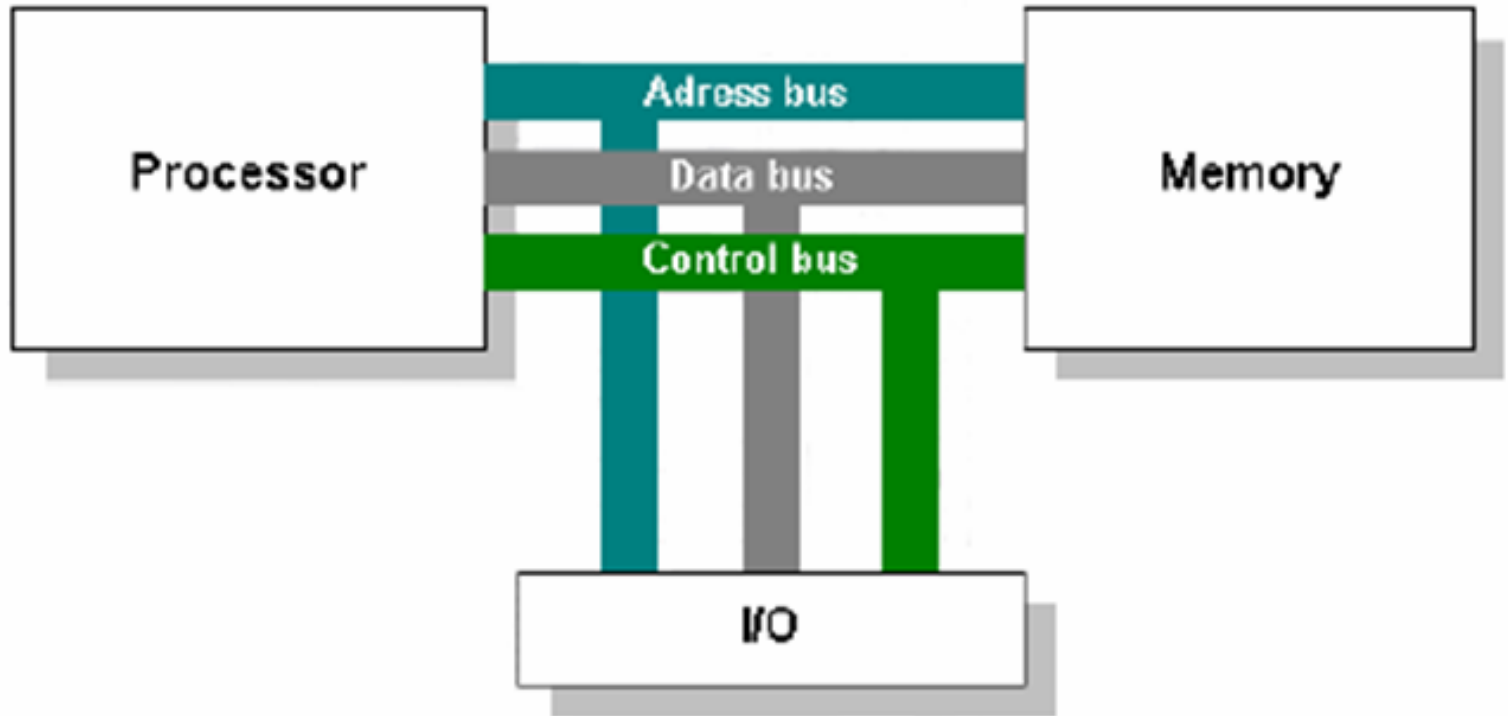
# What is a Bus?

---

- *A communication pathway connecting two or more devices*
  - *Usually broadcast*
  - *Often grouped*
  - *A number of channels in one bus*
- *e.g. 32 bit data bus is 32 separate single bit channels*
- *Power lines may not be shown*



# Specific Buses



# Data Bus

---

- *Carries data*
  - *There is no difference between “data” and “instruction” at this level*
- *Width is a key determinant of performance*
  - *8, 16, 32, 64 bit*

# Address Bus

---

- *Identify the source/ destination of data*
- *e.g. CPU needs to read an instruction (data) from a given location in memory*
- *Bus width determines maximum memory capacity of system*
  - *e.g. MIPS has 32 bit address bus giving 4GB address space*

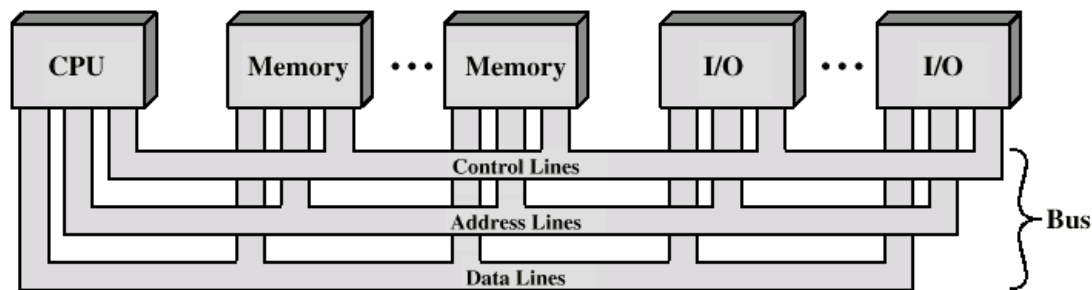
# Control Bus

---

- *Control and timing information*
  - *Memory read/write signal*
  - *Interrupt request*
  - *Clock signals*

# Disadvantage of Buses

- *Creates a communication bottleneck*
  - *Bus bandwidth limits maximum I/O throughput*
- *Maximum bus speed is largely limited by:*
  - *Length of bus*
  - *Number of devices on bus*
  - *Slowest device on bus*



# I/O Transactions

---

- *An I/O Transaction Consists of:*
  - *sending address & command*
  - *sending or receiving data*
- *Types of I/O Transactions:*
  - *Input*
    - *from input devices to memory or processor*
  - *Output*
    - *from memory or processor to output devices*

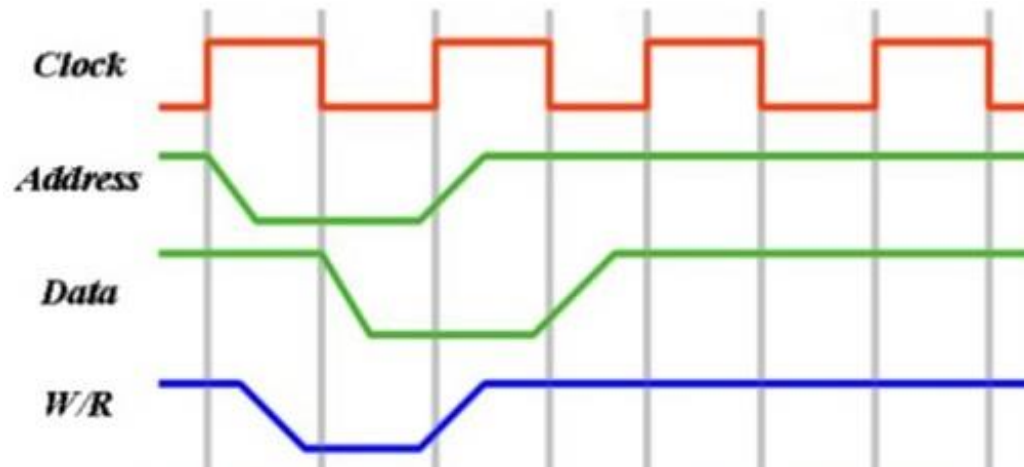
# Sync/Async Interconnection

---

- *Synchronous* interconnection
  - includes a *clock* in the control lines
  - uses a fixed protocol for communicating that is relative to the clock
- *Asynchronous* interconnection
  - not clocked
  - uses a *handshaking* protocol for coordinating usage rather than a clock

# Synchronous Interconnection

- A *clock* in the control lines
- *Fixed* communicating protocol relative to the clock





# Sync Interconnection (pros/cons)

---

- ☺ can be implemented *easily* in a small finite-state machine
- ☺ the bus can run *fast*, and the interface logic will be *small*
- ☺ Every device on the bus must run at the *same clock rate*
- ☹ *cannot be long* if fast (due to clock skew problems)

# Asynchronous Interconnection

---

- *not clocked*
- *uses a **handshaking** protocol for coordinating usage rather than a clock*
  - *consists of a series of steps in which the sender and receiver proceed to the next step only when both parties agree*
  - *The **protocol** is implemented with an additional set of control lines*

## Async Interconnection (pros/cons)

---

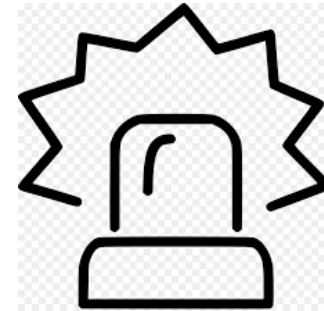
- 😊 can accommodate a *wide variety of devices* of different speeds
- 😊 the bus can be *lengthened* without worrying about clock skew or synchronization problems
- 😞 ?

# Typical Async Protocols

---

## ○ *Strobing*

- *Source-initiated*
- *Destination-initiated*

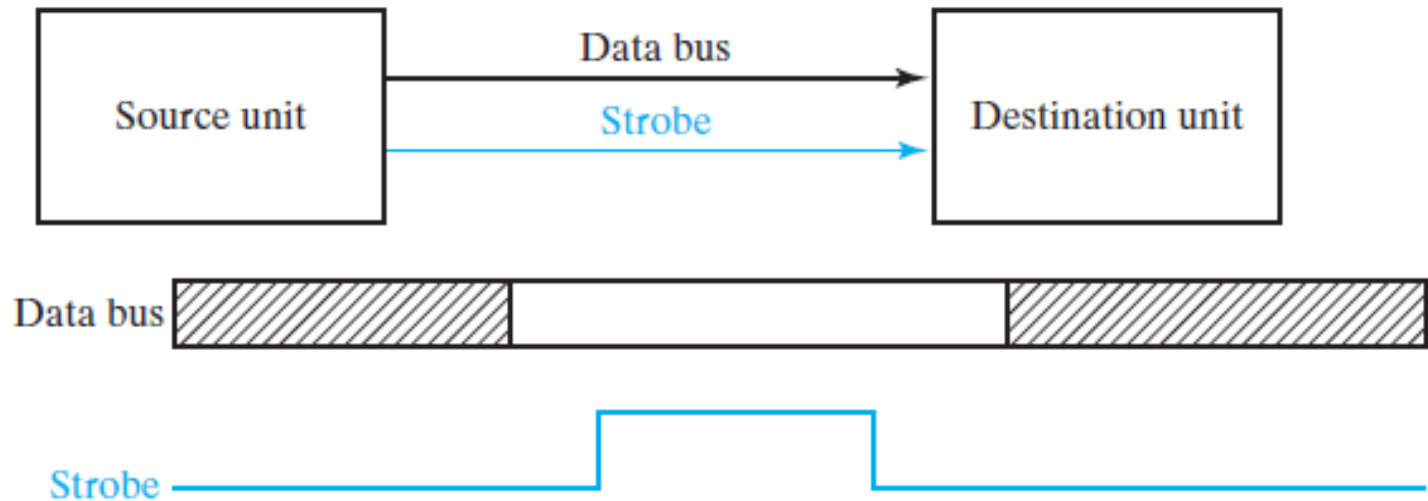


## ○ *Handshaking*

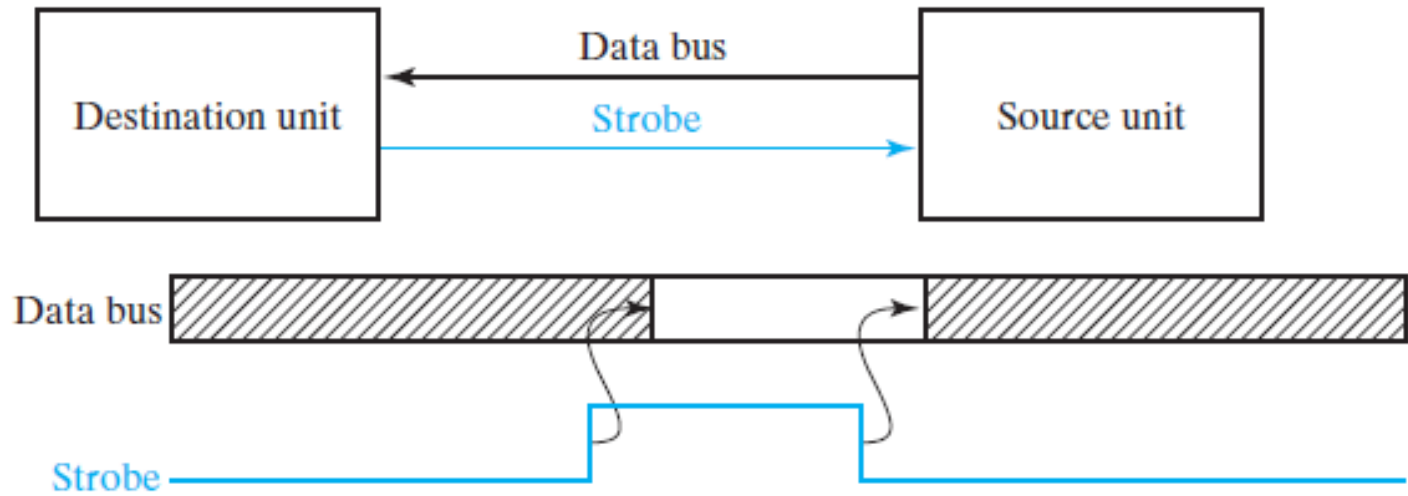
- *Source-initiated*
- *Destination-initiated*



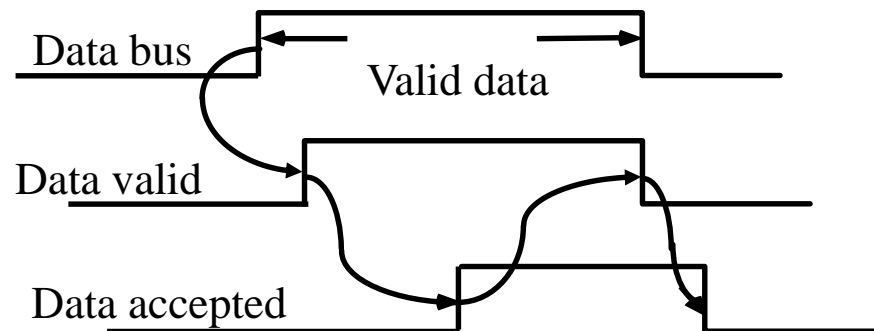
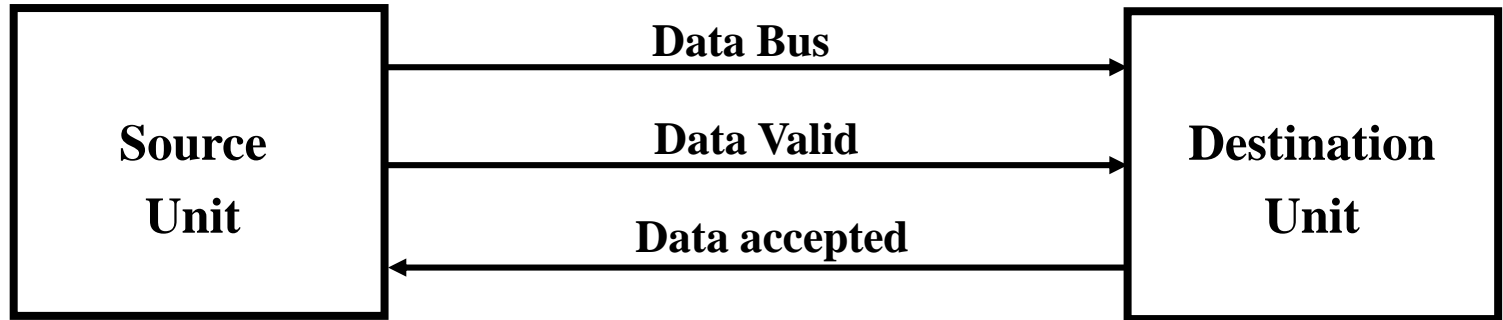
# Source-Initiated Strobe



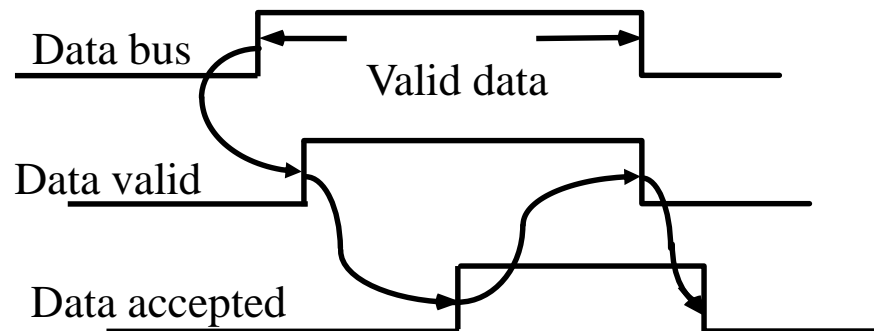
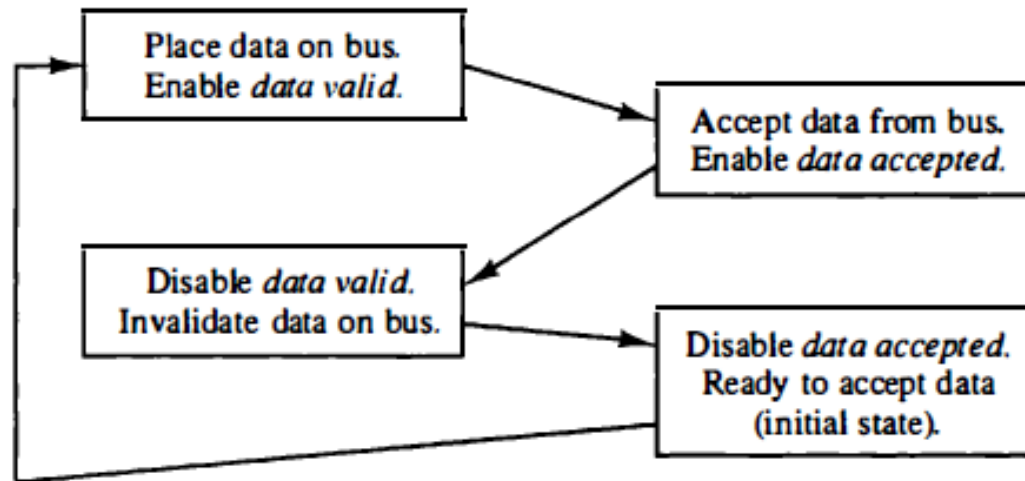
# Destination-Initiated Strobe



## Src-Initiated Transfer via Handshake - 1

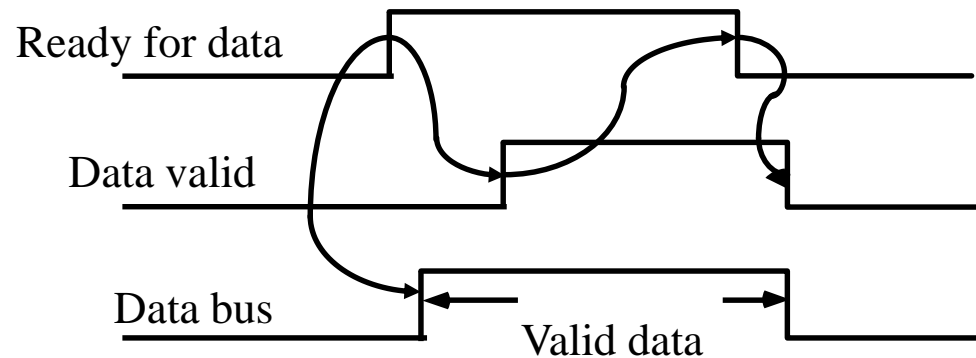
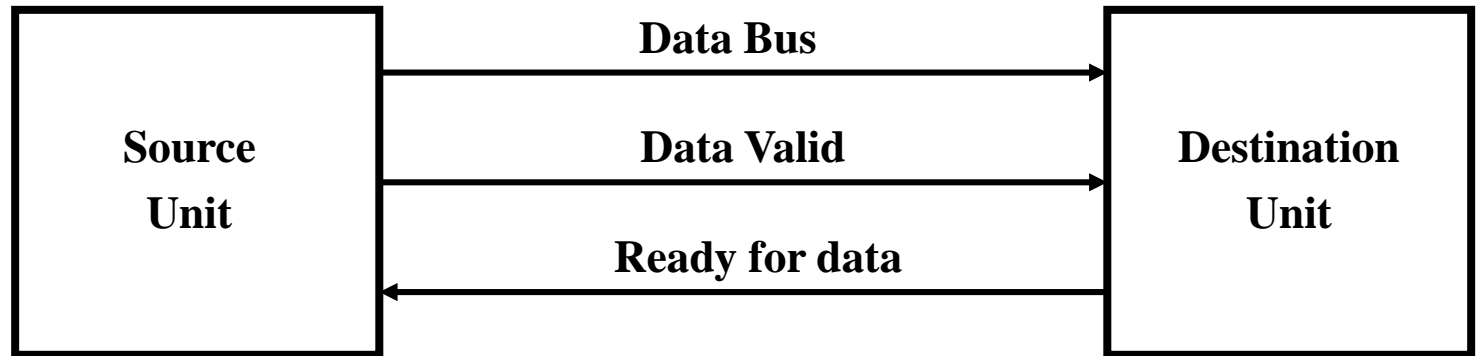


## Src-Initiated Transfer via Handshake - 2

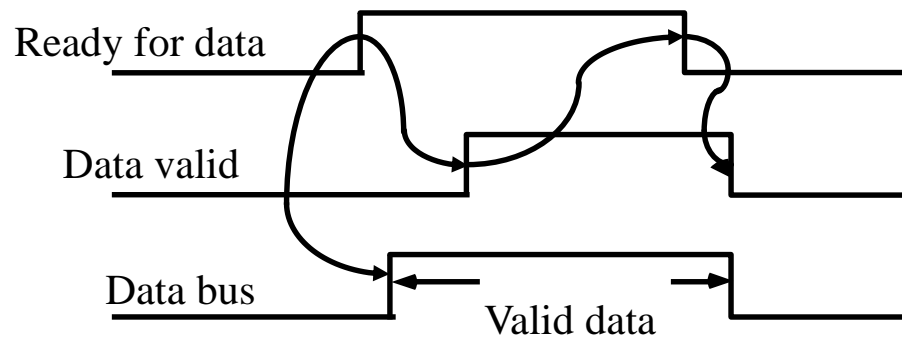
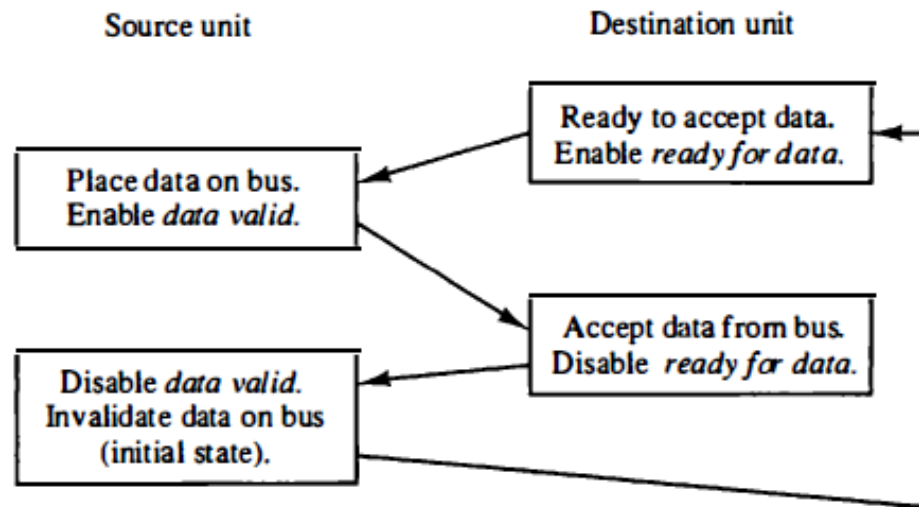




## *Dst-Initiated Transfer via Handshake - 1*



## *Dst-Initiated Transfer via Handshake - 2*



## Still More ...

---

- How is a user I/O request transformed into a device command and communicated to the device?
- What is the role of the *operating system*?

# Operating System?

---

- *OS responsibilities arise from I/O characteristics:*
  - *Multiple programs using the processor **share** the I/O system*
  - *I/O systems often use **interrupts** to communicate information about I/O operations. Interrupts cause a transfer to kernel or supervisor mode, so they must be handled by the OS*
  - *The low-level control of an I/O device is **complex**:*
    - *requires managing a set of concurrent events*
    - *requirements for correct device control are detailed*

# Operating System's Role

---

- *Guarantees that a user's program accesses **only** the portions of an I/O device to which the user has **rights***
  - *e.g., must not allow a program to read or write a file on disk if the owner of the file has not granted access to this program*
  - *In a system with shared I/O devices, protection could not be provided if user programs could perform I/O directly*
- ...

# Operating System's Role (cont.)

---

- Guarantees that a user's program accesses *only* the portions of an I/O device to which the user has *rights*
- Provides *abstractions* for accessing devices by supplying routines that handle low-level device operations
- Handles the *interrupts* generated by I/O devices, just as it handles the *exceptions* generated by a program
- Tries to provide *equitable* access to the shared I/O resources and *schedule* accesses to enhance system throughput

# OS Communication with I/O

---

- *Give commands to the I/O devices:*
  - *read, write, disk seek, ...*
- *Be notified when the I/O device has completed an operation or has encountered an error*
  - *e.g., when a disk completes a seek, it will notify OS*
- *Transfer data between memory and I/O device*
  - *e.g., the block being read from a disk must be moved from disk to memory*

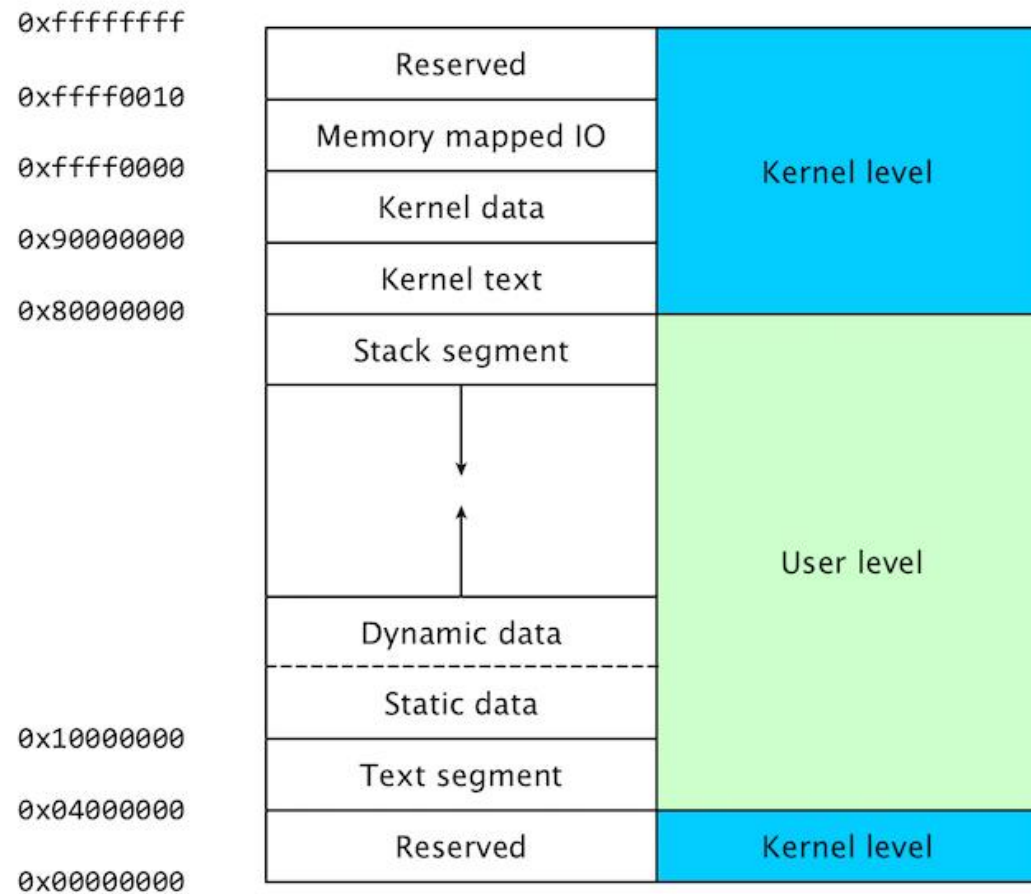
# *Giving Commands to I/O Devices*

---

- *How to Access I/O?*
  - *From ISA perspective*
- *I/O Configuration*
  - *Isolated I/O*
  - *Memory-Mapped I/O*



# Memory-Mapped I/O in MIPS



# Memory-Mapped I/O

---

- Performed by simple *load/store* instructions
- A subset of unused memory addresses mapped to registers of external devices
- Memory & devices on bus programmed to respond only to their own address ranges

# Memory-Mapped I/O

---

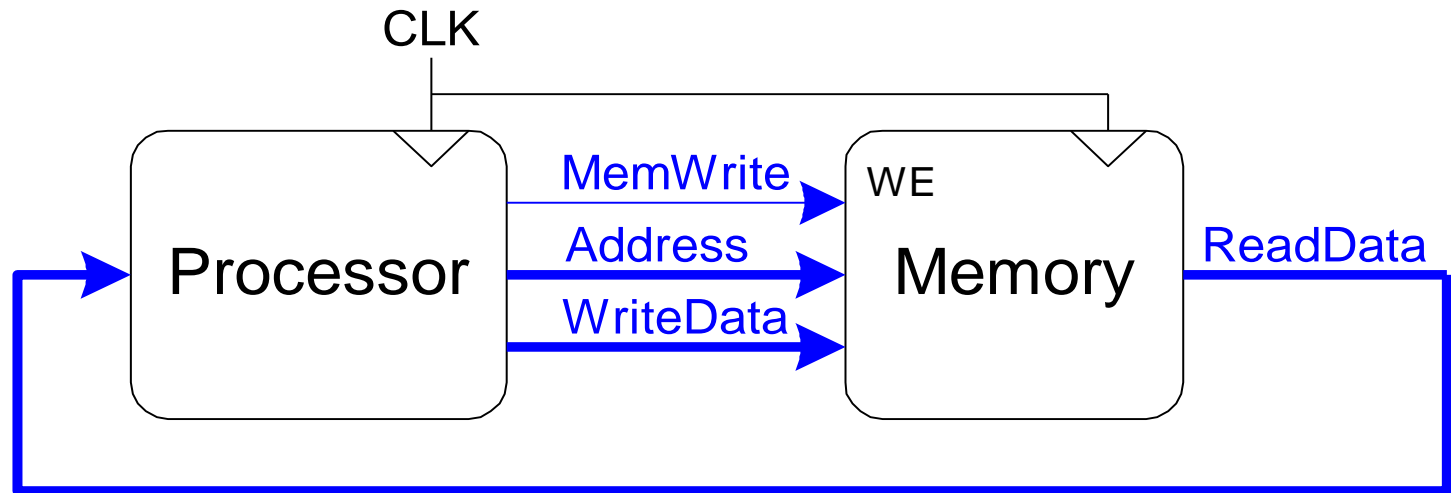
## 😊 Pros

- *Easy to handle memory locations (software perspective)*
- *Wide range of addressing modes*

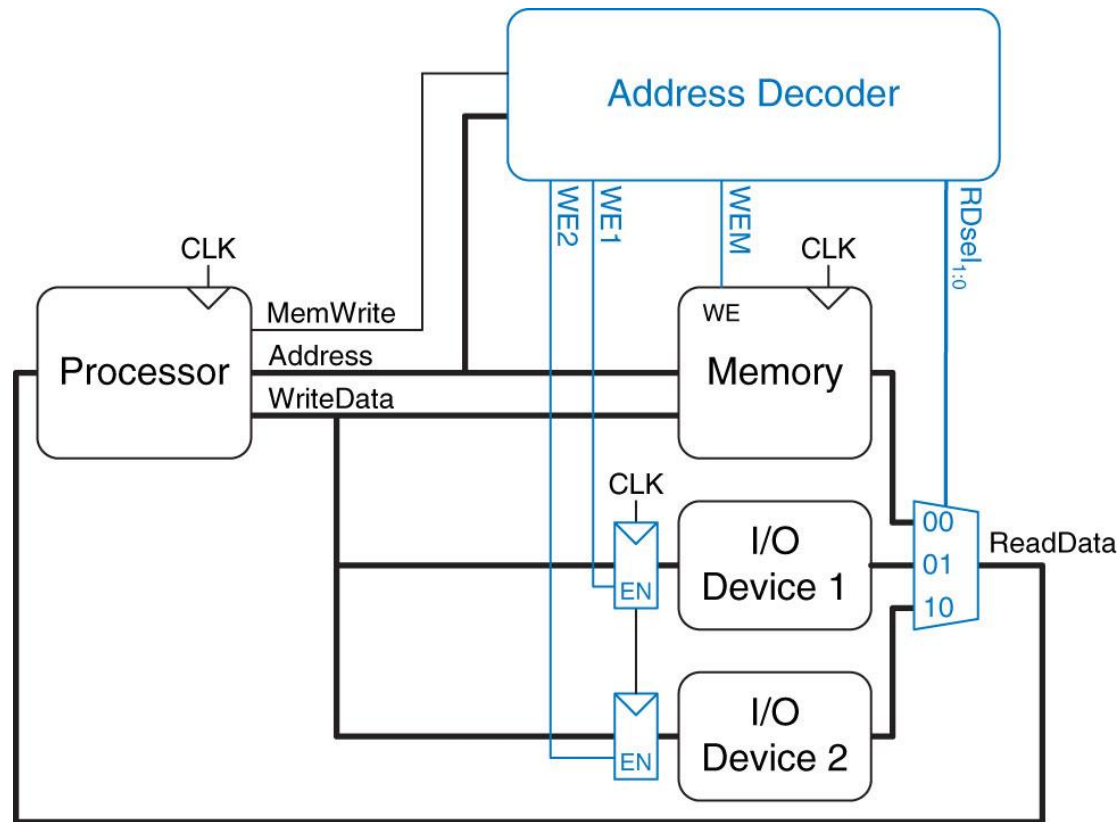
## 😞 Cons

- *Caching may cause to memory incoherency issue*
- *Slow (consumes CPU cycles)*

# Memory Interface

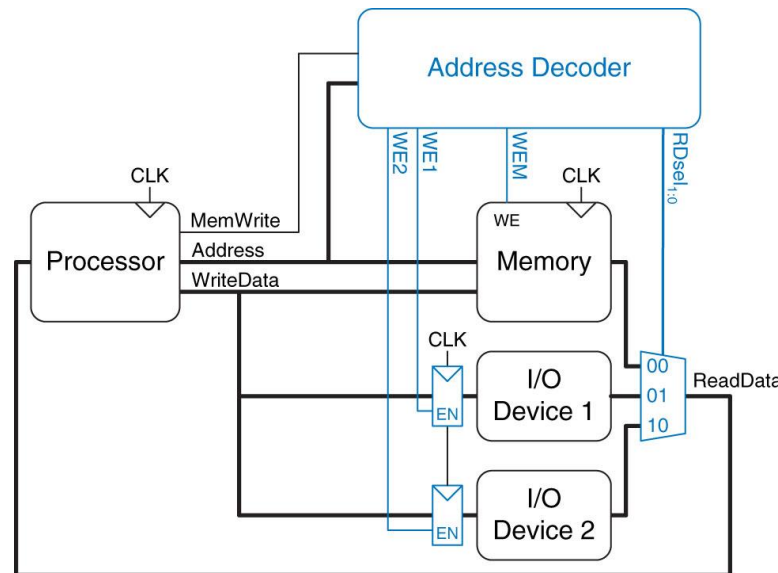


# Memory-Mapped I/O Hardware



# Memory-Mapped I/O Example

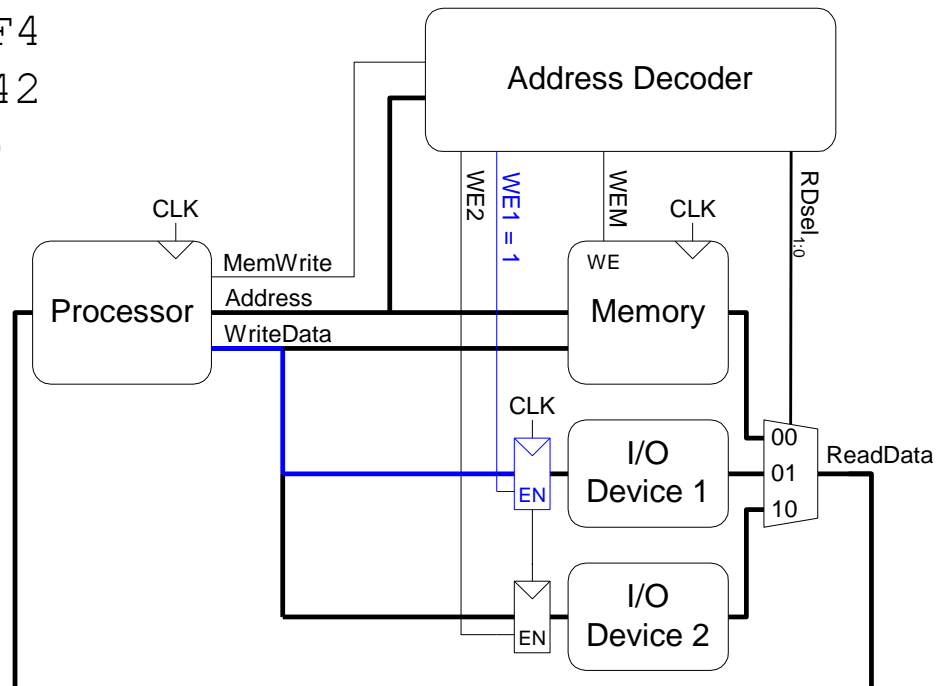
- I/O Device 1 is assigned the address 0xFFFFFFFF4
  - Write the value 42 to I/O Device 1
  - Read value from I/O Device 1 and place in r1



# Memory-Mapped I/O Example

- Write the value 42 to I/O Device 1 (0xFFFFFFFF4)

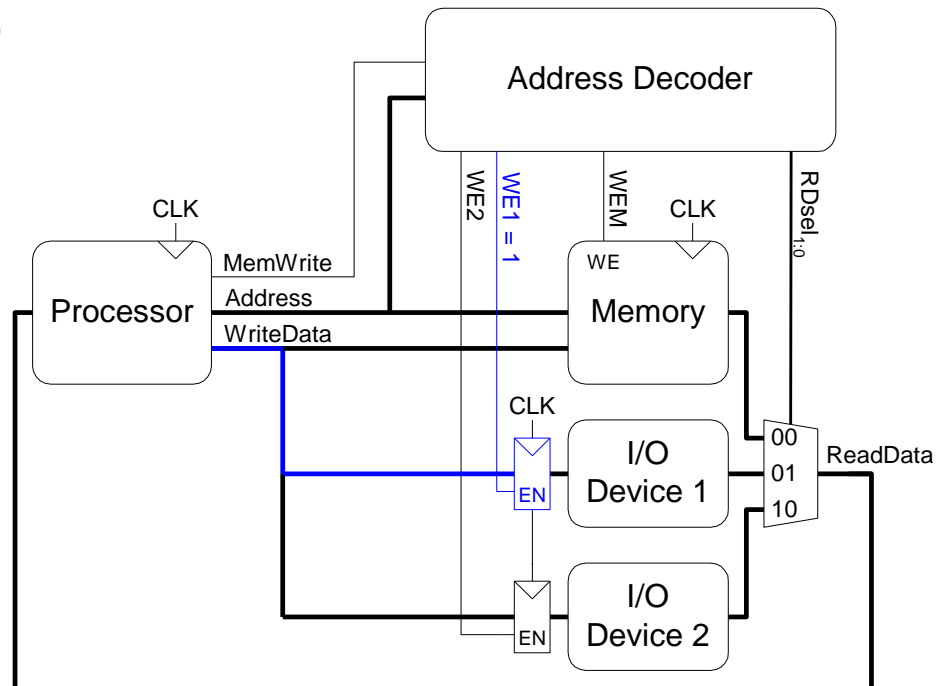
```
lui $t1, 0xFFFF
ori $t1, 0xFFF4
addi $t0, $0, 42
sw $t0, 0($t1)
```



# Memory-Mapped I/O Example

- Read value from I/O Device 1 and place in \$t0

```
lw $t0, 0($t1)
```





# Isolated I/O

---

- *Dedicated I/O instructions*
  - *Part of ISA*
- *Output*
  - *values are written to an output register*
  - *on the output pins of an external port*
- *Input*
  - *values are read from an input register*
  - *From the input pins of an external port*
- *Example: Intel x86 IN/OUT instructions*

## *I/O-Mapped I/O (cont.)*

---

### 😊 *Pros*

- *Low latency*
  - *Simpler decoding*

### 😞 *Cons*

- *Not implemented by all CPUs*
- *Limited addressing*
  - *Predetermined # of I/O signals*

# OS Communication with I/O

---

- Give commands to the I/O devices:
  - read, write, disk seek, ...
- Be notified when the I/O device has completed an operation or has encountered an error
  - e.g., when a disk completes a seek, it will notify OS
- Transfer data between memory and I/O device
  - e.g., the block being read from a disk must be moved from disk to memory

# Programmed I/O vs. Interrupt

---

## ○ Imagine

- scrolling your mouse in screen
- typing by your keyboard
- accessing *slow*-running I/O devices

## ○ How the processor responds?

- *Programmed I/O (aka. Polling)*
- *Interrupt-driven*

# Programmed I/O

---

- *Checks the status of I/O device periodically to determine the need to service*
  - 😊 *Simple*
  - 😊 *More predictable I/O overhead*
  - 😞 *can waste a lot of processor time*
- *Used in real-time applications because of predictability*

# Interrupt-Driven I/O

- A scheme that *interrupts* the processor to indicate that an I/O device needs *attention*



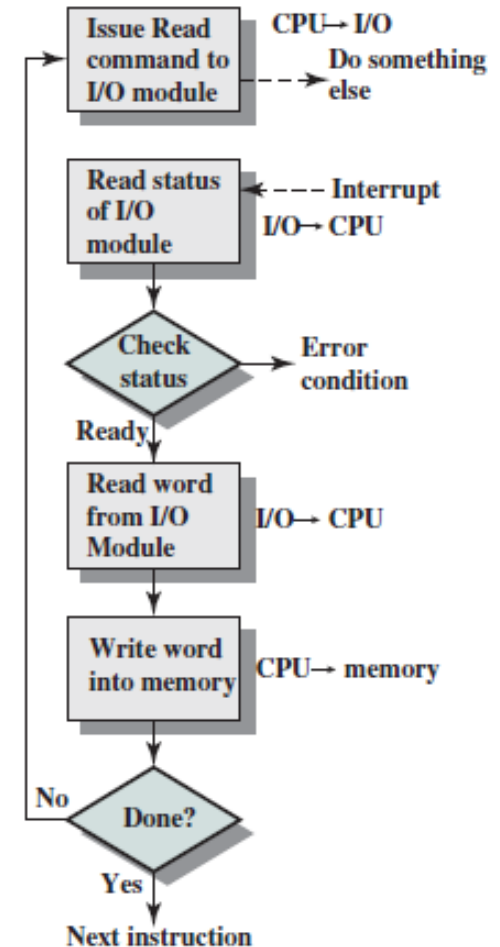
# Interrupt-Driven I/O

---

- A scheme that *interrupts* the processor to indicate that an I/O device needs *attention*
  - ☺ Overcomes CPU waiting
  - ☺ No repeated CPU checking of device
  - ☺ I/O device interrupts when ready
- I/O device
  - *Raises* the interrupt
  - *Identifies* itself
- The OS gets control and responds to the interrupt

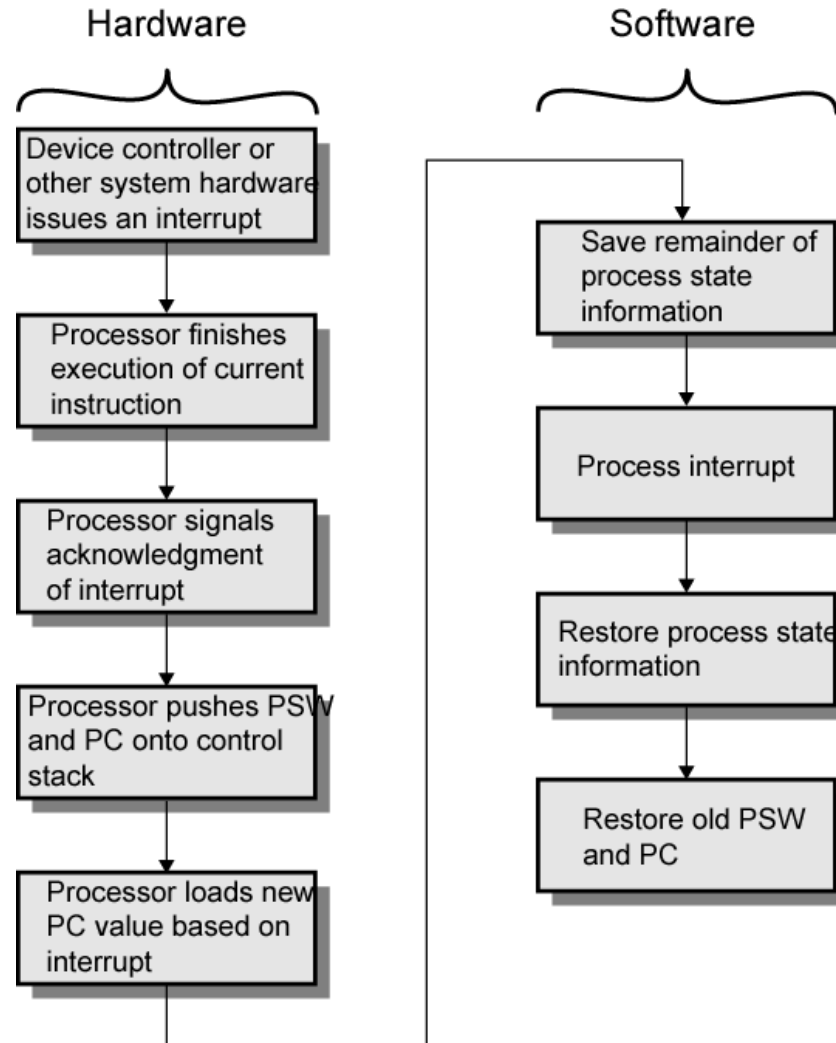
# Interrupt Driven I/O Basic Operation

- CPU issues read command
- I/O module gets data from peripheral whilst CPU does other work
- I/O module interrupts CPU
- CPU requests data
- I/O module transfers data





# Simple Interrupt Processing

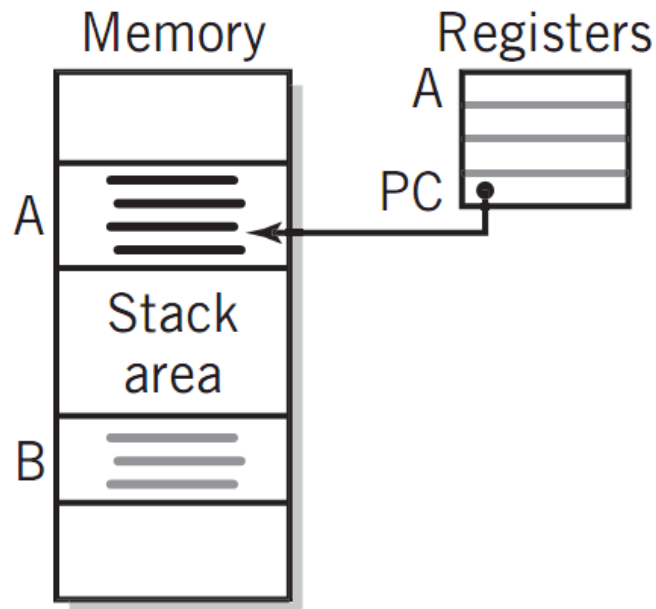


# CPU Viewpoint

---

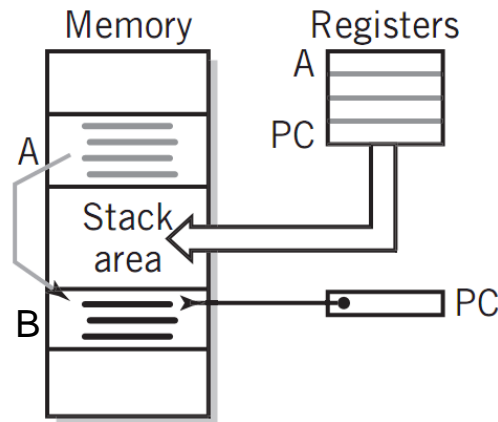
- *Issue read command*
- *Do other work*
- *Check for interrupt at end of each instruction cycle*
- *If interrupted:*
  - *Save context (registers)*
  - *Process interrupt*
    - *Fetch data & store*

# Servicing Interrupts - 1



*Before interrupt arrives, program A is executing  
The PC (program counter) points to the current instruction*

# Servicing Interrupts - 2

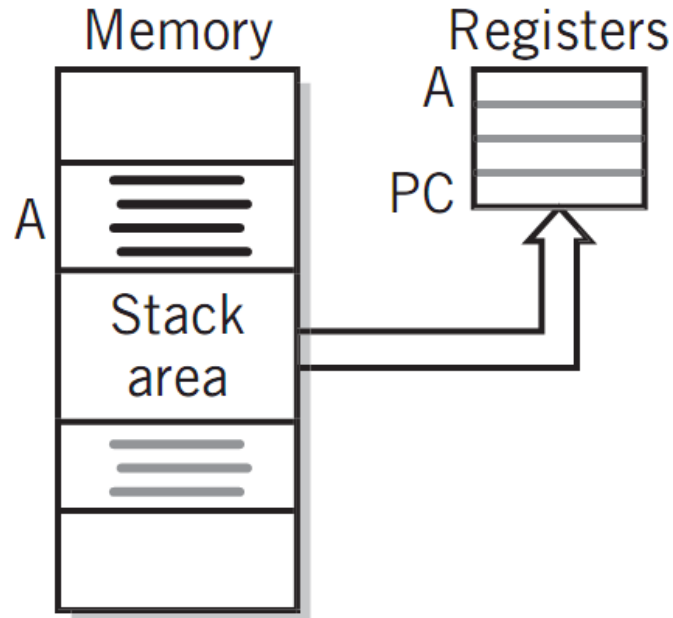


When CPU receives the interrupt, the *current instruction is completed* and all the *registers are saved* in the stack area (or in a special area known as a process control block)

The *PC is loaded* with the starting location of program B, the interrupt handler program

This causes a jump to B, which becomes the executing program

# Servicing Interrupts - 3



*When the interrupt routine is complete, the registers are restored, including the PC and the original program resumes exactly where it left off*

# *Design Issues*

---

- *How do you identify the module issuing the interrupt?*
- *How do you deal with multiple interrupts?*
  - *i.e. an interrupt handler being interrupted*

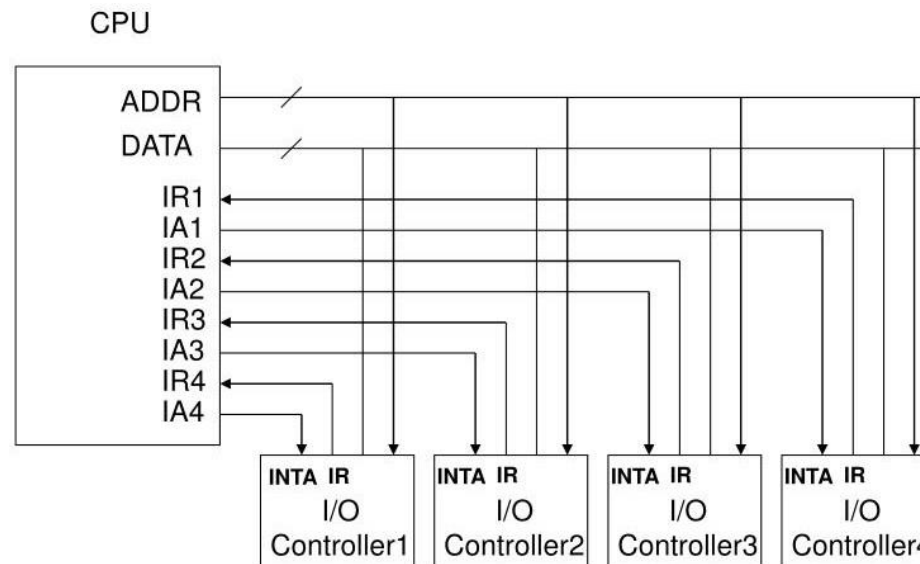
# Identifying Interrupting Module

---

- *Multiple interrupt lines*
- *Software poll*
- *Daisy Chain or Hardware poll*
- *Parallel Priority Hardware*
- *Bus arbitration*

# Multiple Interrupt Lines

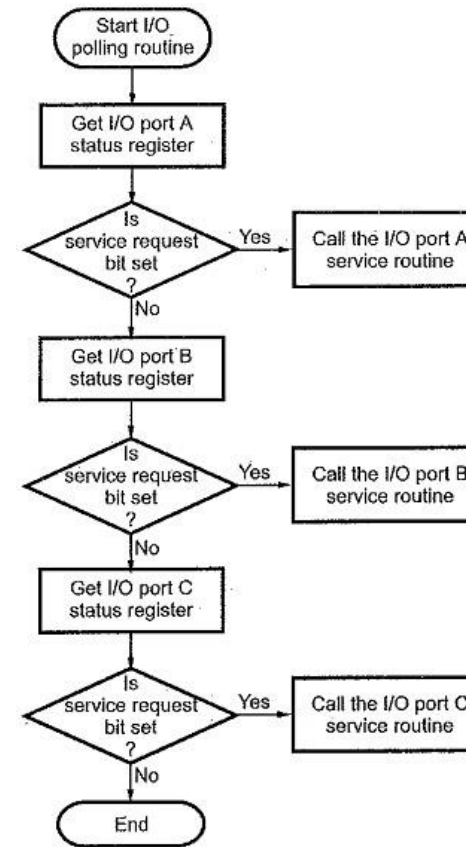
- Different line for each module
- Limits number of devices





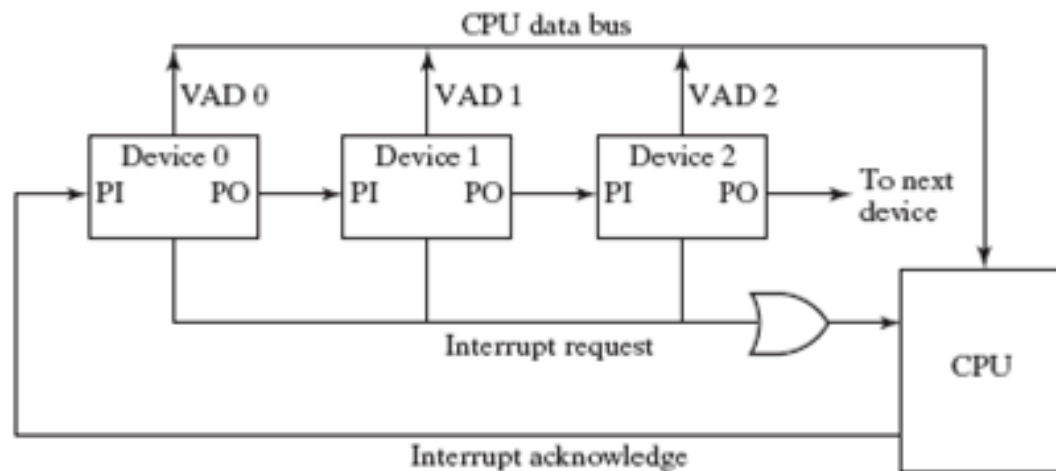
# Software Poll

- CPU asks each module in turn
- Time consuming

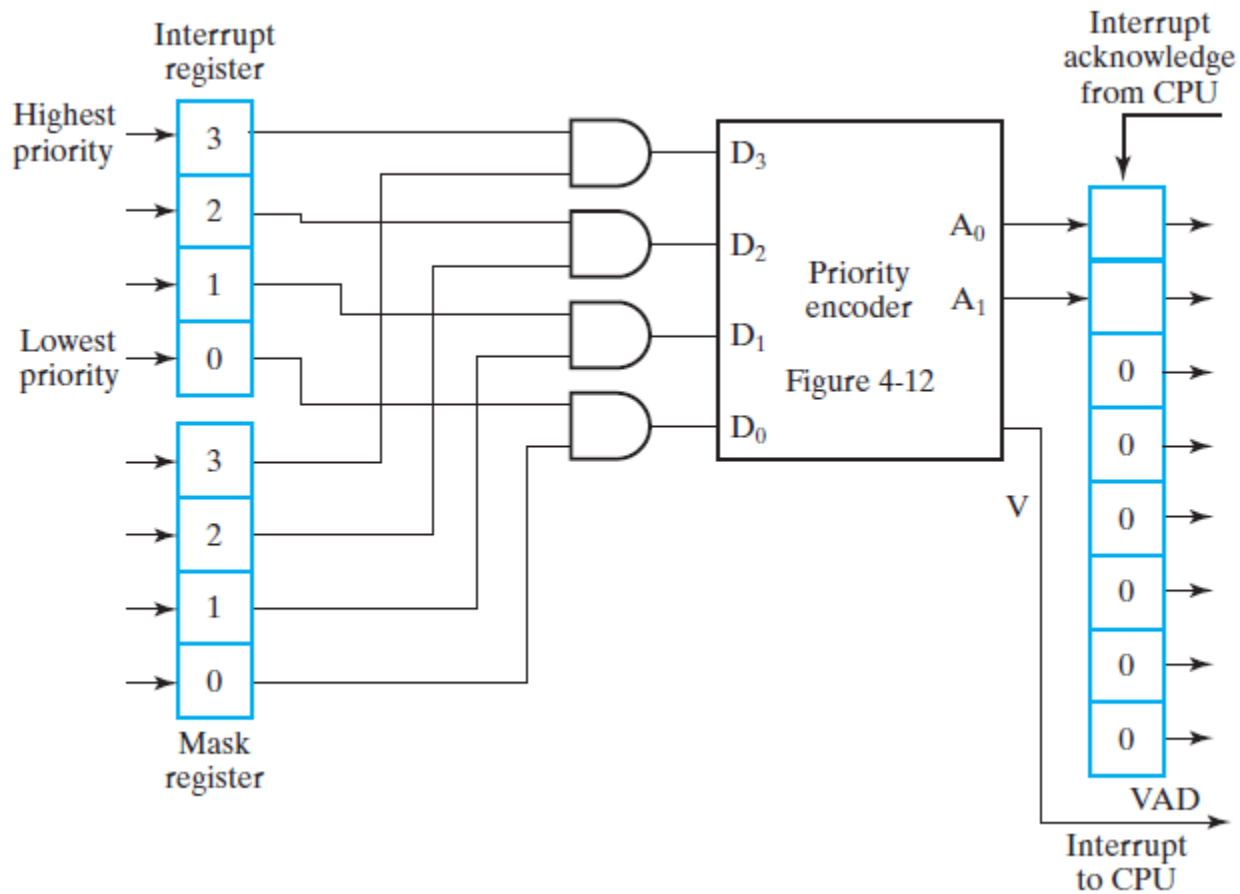


# Hardware Poll (Daisy Chain)

- Interrupt Acknowledge sent down a chain
- Module responsible places **vector** on bus
- CPU uses vector to identify interrupt

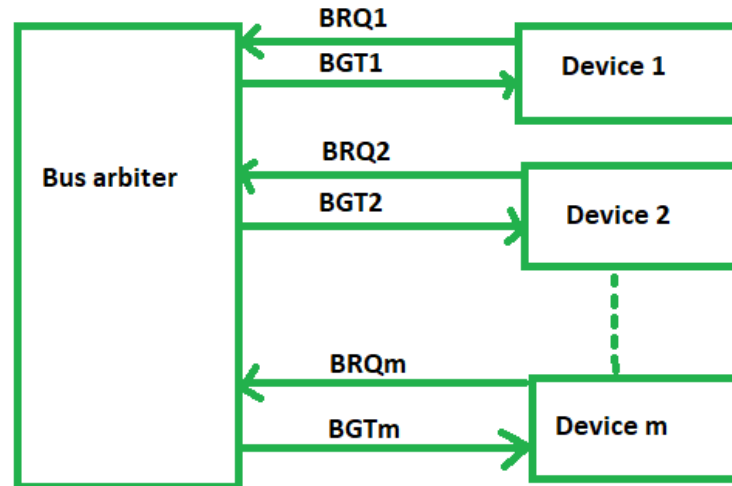


# Parallel Priority Hardware



# Bus Arbitration

- *Module claims the bus before it can raise interrupt*
- *The processor detects the interrupt and responds on the interrupt acknowledge line*
- *The requesting module places its vector on data lines*



# Multiple Interrupts

- Each interrupt has a **priority**
- Higher priority lines can interrupt lower priority lines

Non-Maskable Interrupts (NMI)



High Priority Interrupts



Medium Priority Interrupts



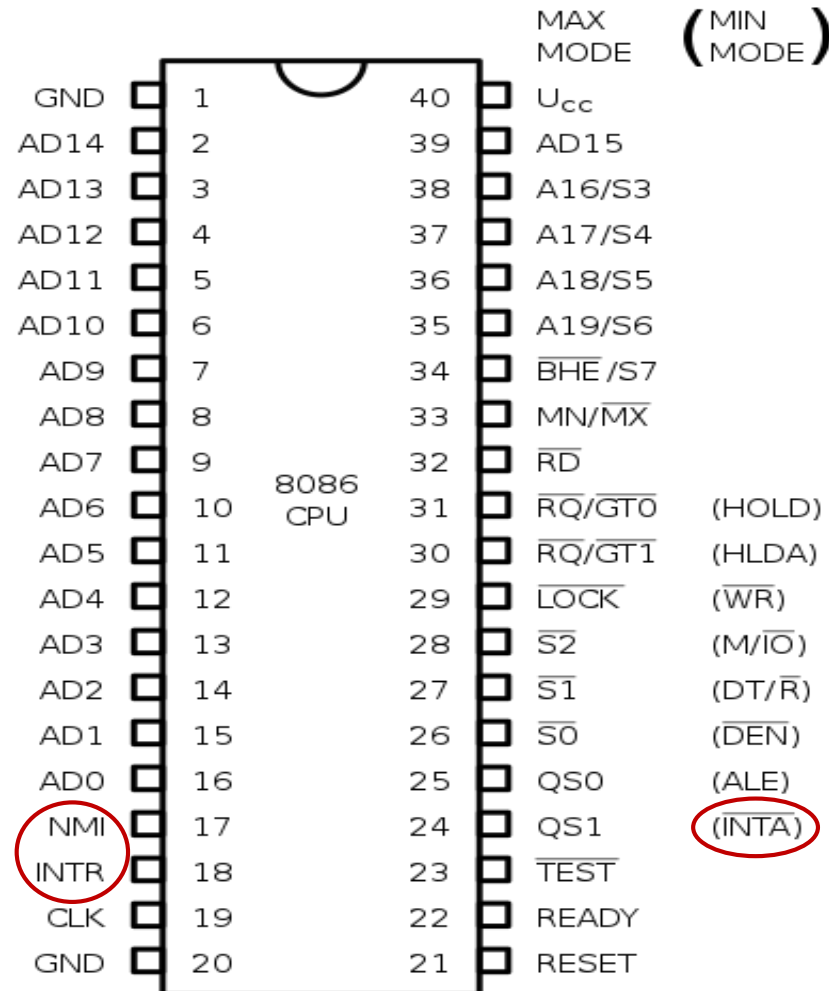
Low Priority Interrupts

# Handling Priority

---

- *Multiple interrupt lines*
  - *processor picks the interrupt line with highest priority*
- *Software poll*
  - *order in which modules are polled determines priority*
- *Daisy Chain or Hardware poll*
  - *order of modules on a daisy chain determines priority*
- *Parallel Hardware priority*
  - *Use of priority encoder*
- *Bus arbitration*
  - *only current bus master can interrupt*

# 8086 Chip



# OS Communication with I/O

---

- Give commands to the I/O devices:
  - read, write, disk seek, ...
- Be notified when the I/O device has completed an operation or has encountered an error
  - e.g., when a disk completes a seek, it will notify OS
- Transfer data between memory and I/O device
  - e.g., the block being read from a disk must be moved from disk to memory



# High-bandwidth I/O

- Programmed I/O & interrupt are suitable for *low* bandwidth devices
  - Data transfer is done in small no of bytes
- For *higher* bandwidth (e.g. hard disks)
  - Programmed I/O is used in real-time applications
  - Interrupt driven approach
    - Helps OS to work on other tasks while actual I/O transfer is done
    - But the overhead of each transfer is still intolerable
- Solution: *Direct Memory Access (DMA)*

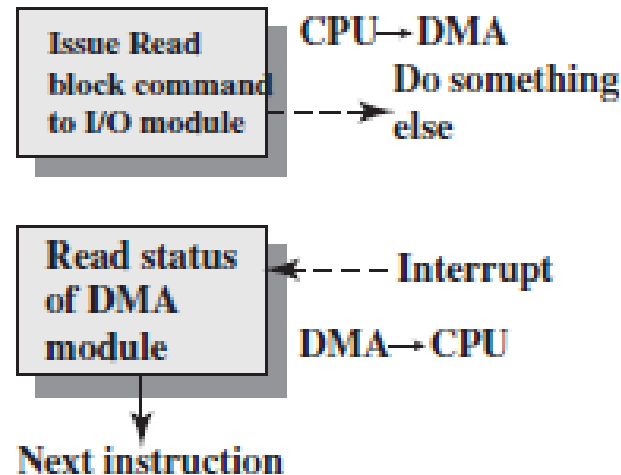
# Direct Memory Access (DMA)

---

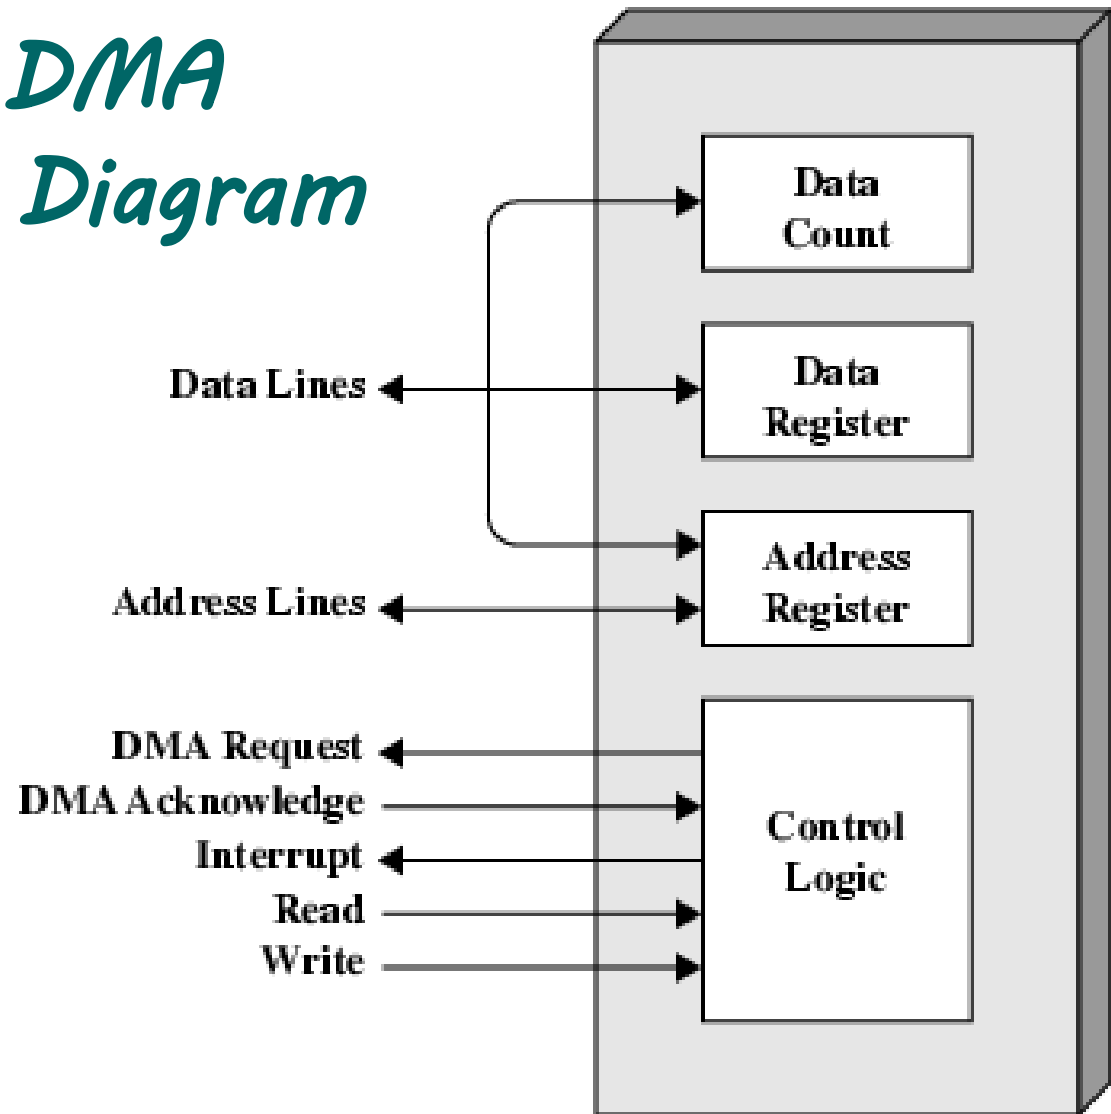
- A mechanism that provides a device controller with the ability to **transfer data directly** to or from the **memory** without involving the processor
- Implemented with a specialized controller that transfers data between an I/O device and memory **independent of the processor**

# DMA Controller

- Additional Module (hardware) on bus
- Becomes **bus master** and directs reads or writes between itself and memory



# Typical DMA Module Diagram

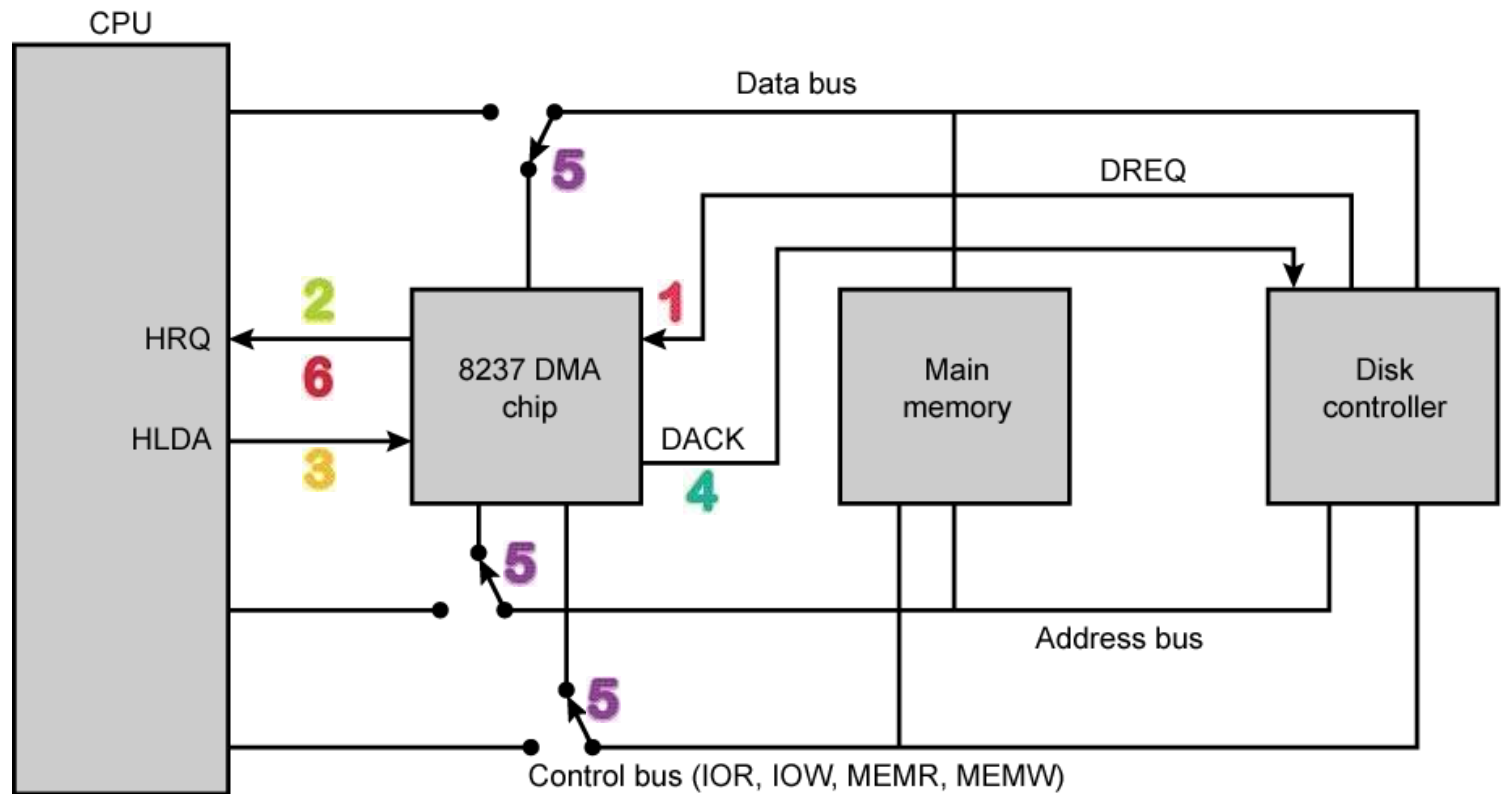


# DMA Operation

---

- CPU tells DMA controller
  - Read/Write
  - Device address
  - Starting address of memory block for data
  - Amount of data to be transferred
- CPU carries on with other work
- DMA controller deals with transfer
- DMA controller sends interrupt when finished

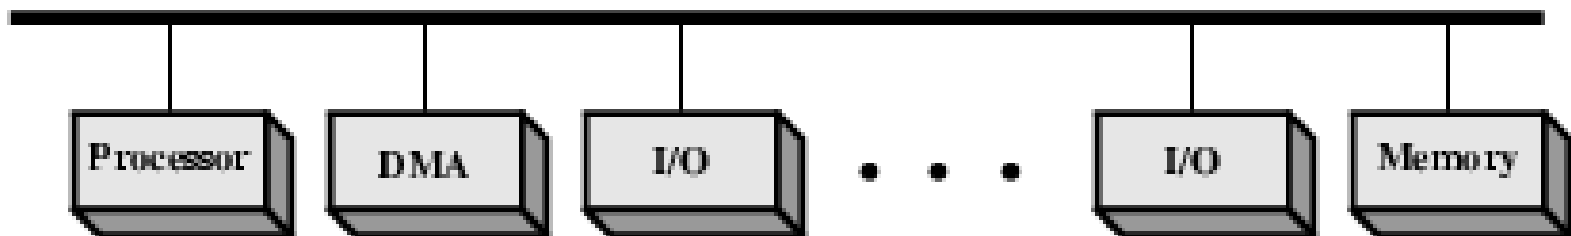
## DMA Usage of Systems Bus



DACK = DMA acknowledge  
DREQ = DMA request  
HLDA = HOLD acknowledge  
HRQ = HOLD request

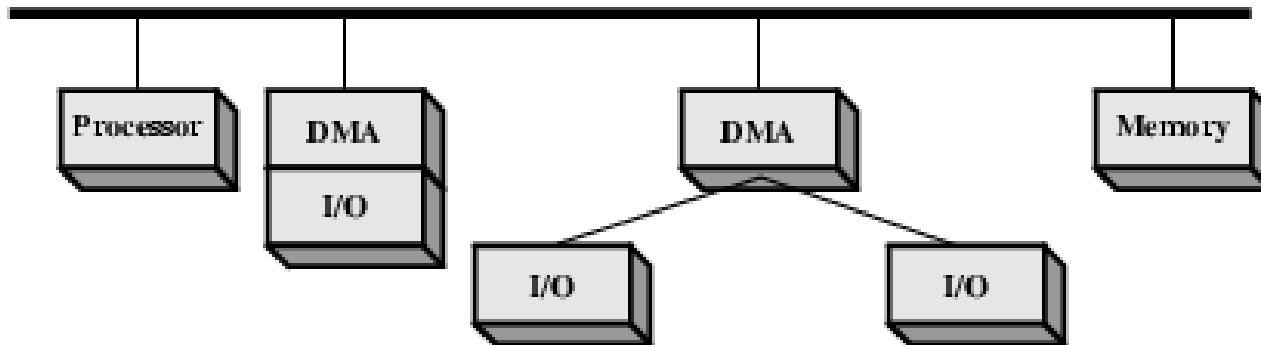
# DMA Configurations (1)

- Single Bus, Detached DMA controller
- Each transfer uses bus twice
  - I/O to DMA then DMA to memory
- CPU is suspended twice



## DMA Configurations (2)

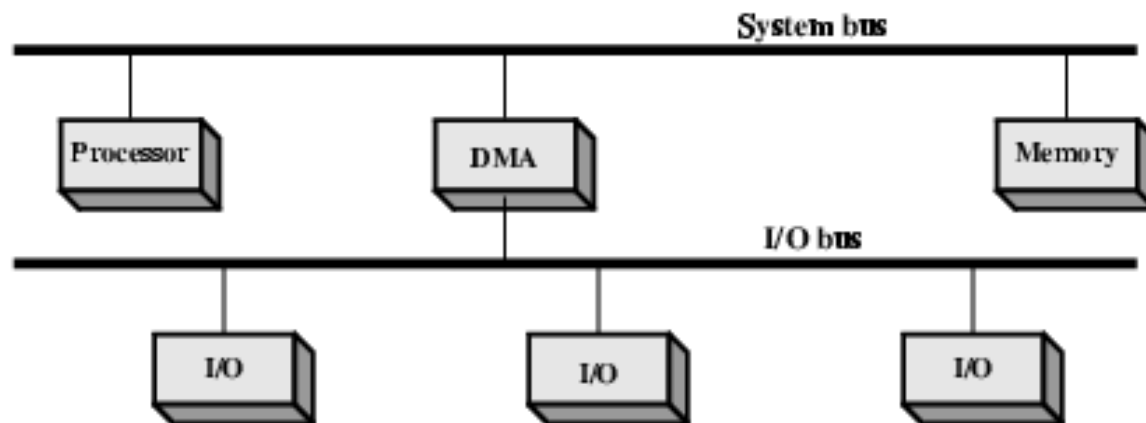
- *Single Bus, Integrated DMA controller*
- *Controller may support >1 device*
- *Each transfer uses bus once*
  - *DMA to memory*
- *CPU is suspended once*





## DMA Configurations (3)

- *Separate I/O Bus*
- *Bus supports all DMA enabled devices*
- *Each transfer uses bus once*
  - *DMA to memory*
- *CPU is suspended once*



# Review

---

- *I/O devices (Instances & Characteristics)*
- *I/O transactions*
  - *Sync & Async data transfer*
  - *Handshaking*
- *ISA perspective*
  - *Memory-mapped vs. isolated I/O*
  - *Interrupt vs. Programmed I/O*
  - *Direct Memory Access (DMA)*

